

Replication for Multiple Partition Keys



Leonard Lobel

CTO, SLEEK TECHNOLOGIES

lennilobel.wordpress.com



Curing Analysis Paralysis

Choose a partition key

Based on your most common querying patterns

Only ONE partition key

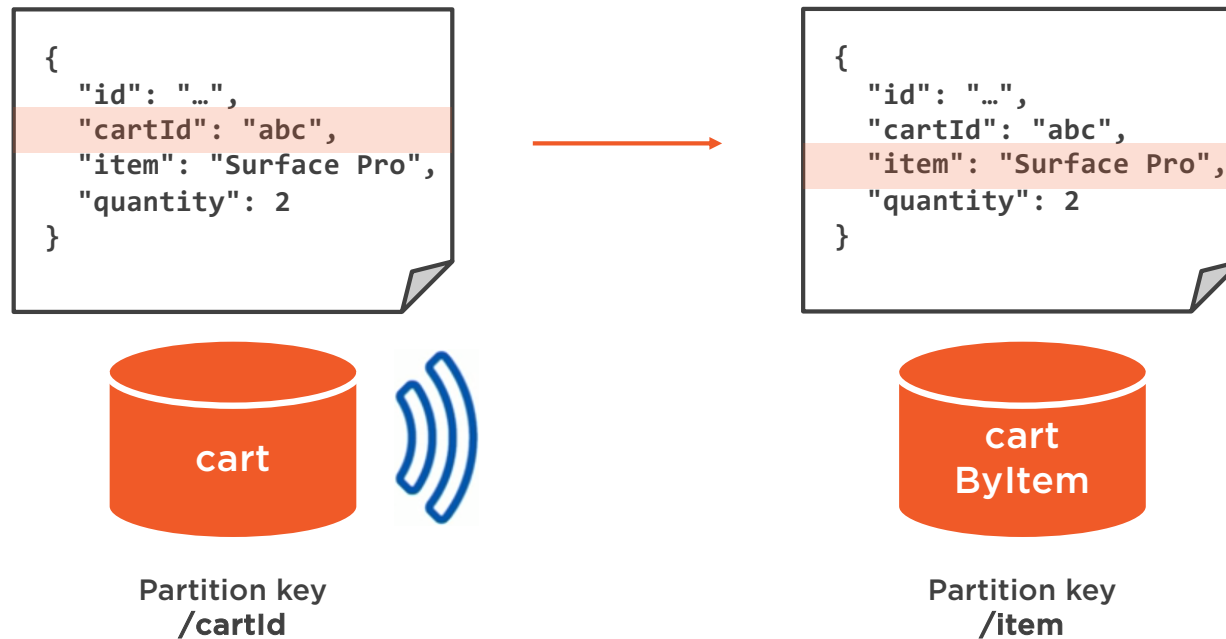
What if you have more than just one common querying pattern?

Replication

Synchronize multiple containers with the same data, partitioned differently



Using Change Feed for Replication



```
SELECT * FROM c WHERE c.cartId = 'abc'
```

```
SELECT * FROM c WHERE c.item = 'Surface Pro'
```

Partition key



Processing Updates and Deletes

Updates

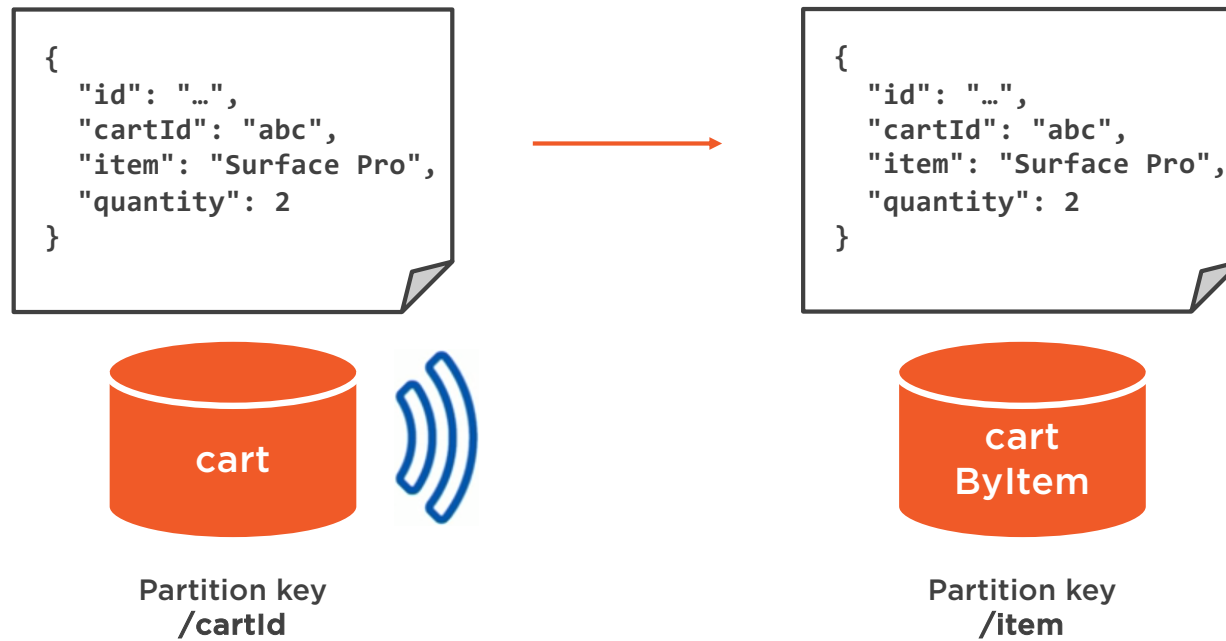
Interim updates are not retained by the change feed

Deletes

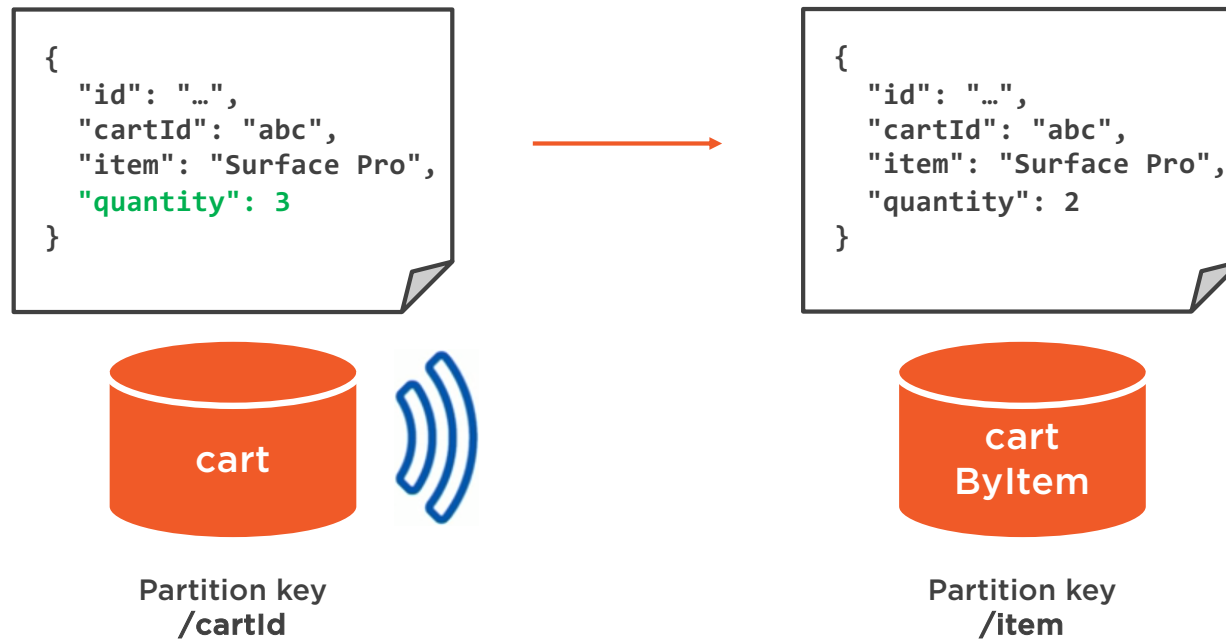
Deletes are not written to the change feed



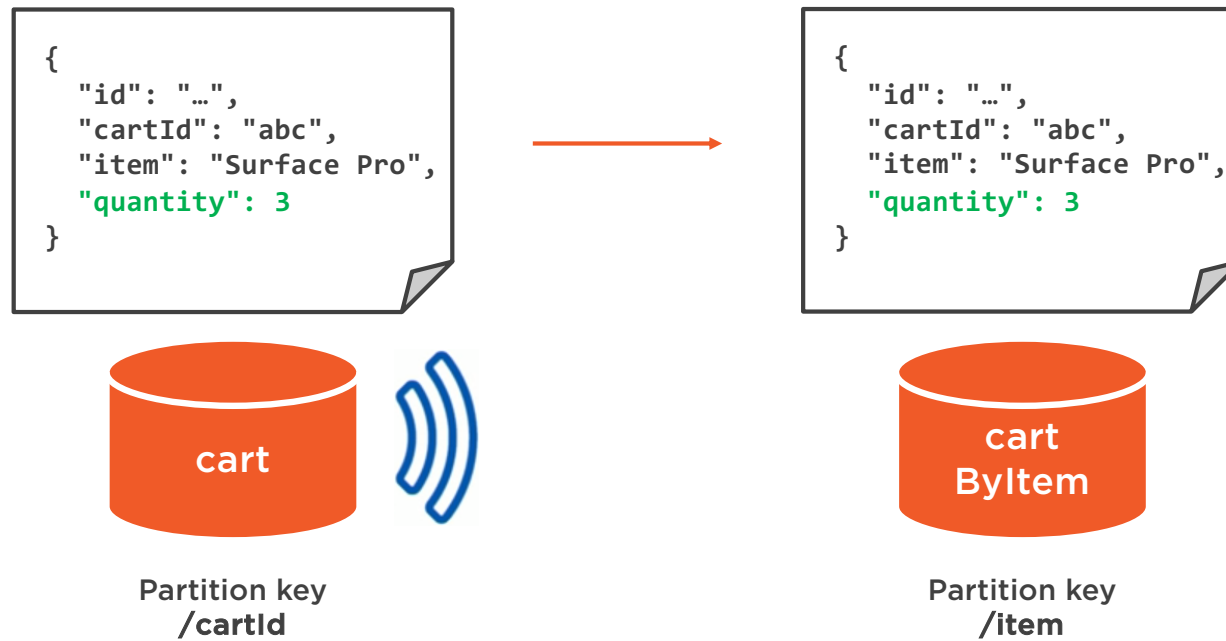
Processing Updates and Deletes



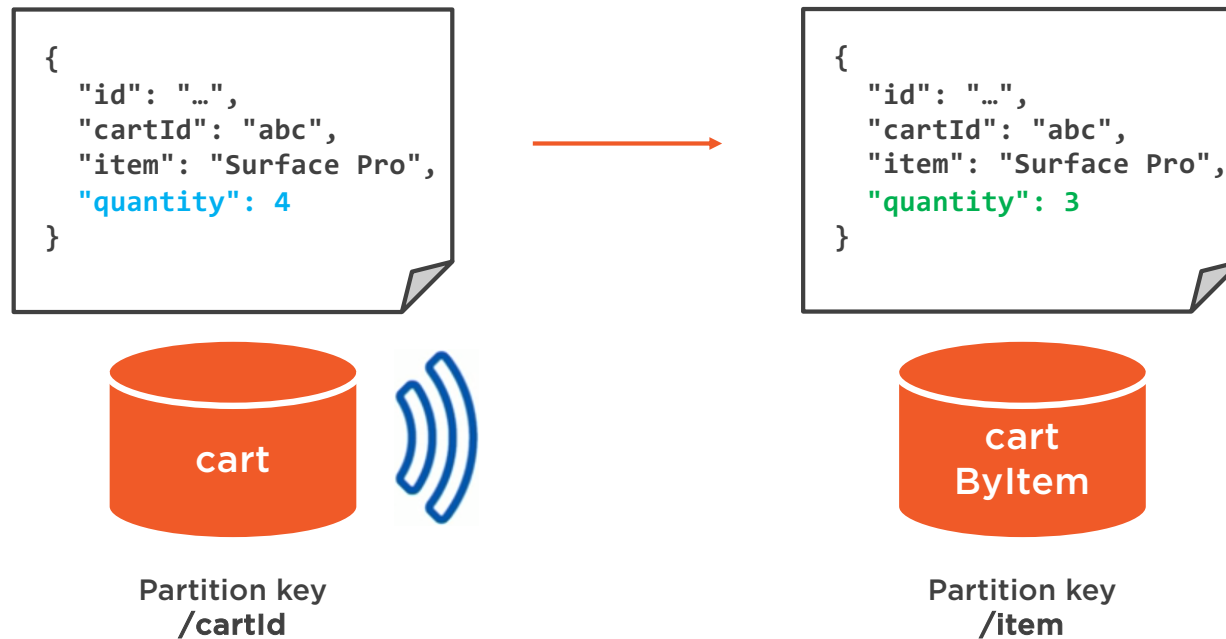
Processing Updates and Deletes



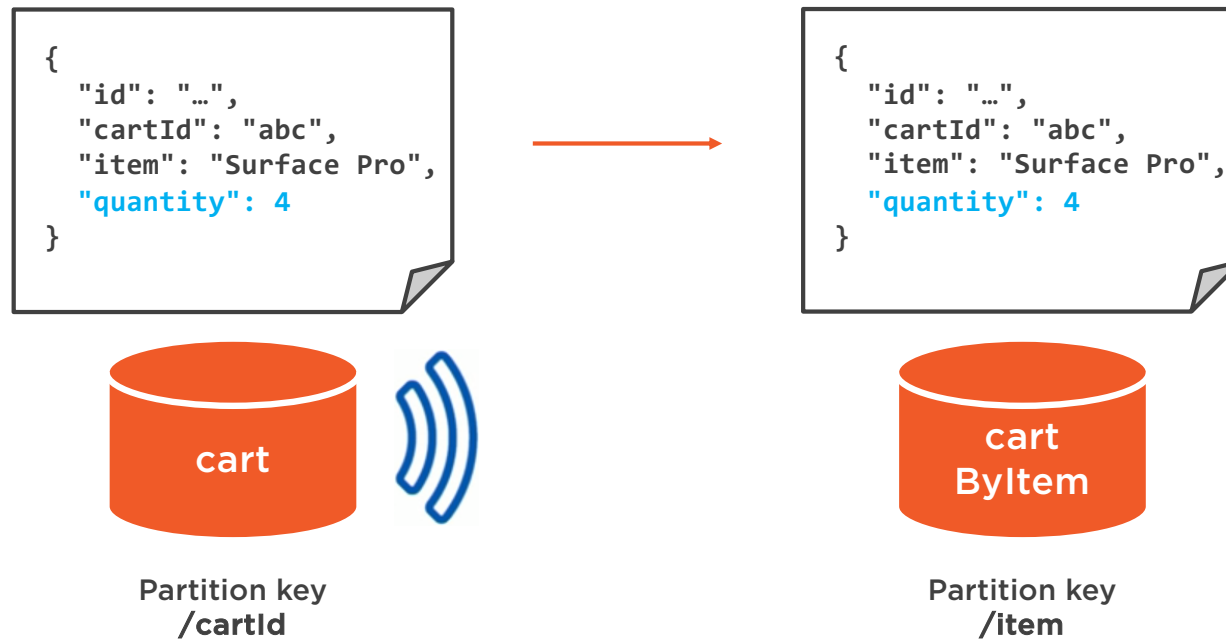
Processing Updates and Deletes



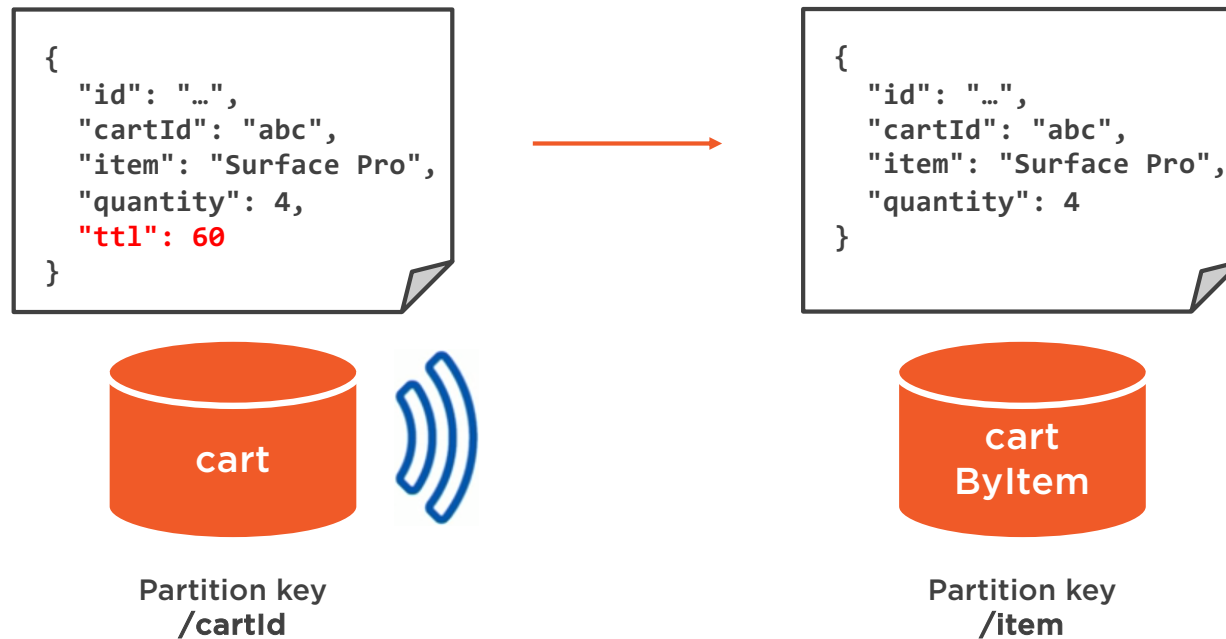
Processing Updates and Deletes



Processing Updates and Deletes



Processing Updates and Deletes



Using Azure Functions

CFP Library

Host is deployed to an
Azure service

E.g., web job, worker role

Dedicated capacity

Azure Functions

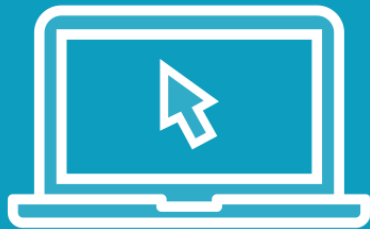
Serverless, reactive
functions

Shared capacity

Dedicated capacity with
AF Premium



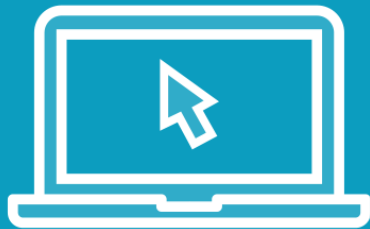
Demo



Preparing for Replication



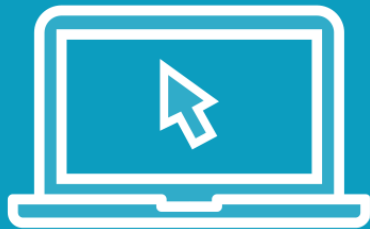
Demo



Creating the Replication Microservice



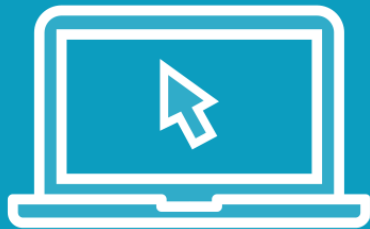
Demo



Testing the Replication Microservice



Demo



Deploying the Replication Microservice



Summary



Replication microservice

- Multiple containers with identical data
- Each partitioned to satisfy different query patterns
- Recommended only for read-heavy workloads

Updates and deletes

- Interim updates are dropped
- Deletes are dropped (use TTL)

Azure Functions

- CFP Library wrapper
- Serverless, reactive functions
- Create, test, and deploy a microservice

