

Denormalizing the Non-relational Data Model



Leonard Lobel

CTO, SLEEK TECHNOLOGIES

lennilobel.wordpress.com



Handling Relational Workloads

Cosmos DB is non-relational

No JOINS
No relational constraints

Paradigm shift

Need to use different techniques to implement relations between entities

Denormalize

Duplicate specific properties
Satisfy queries with a single request



Querying Products

```
{
  "id": "...",
  "categoryId": "...",
  "sku": "...",
  "name": "...",
  "price": "...",
  "tagIds": [
    "...",
    "..."
  ]
}
```

Category name?

Tag names?

product

```
SELECT * FROM c WHERE c.categoryId = '<categoryId>'
```

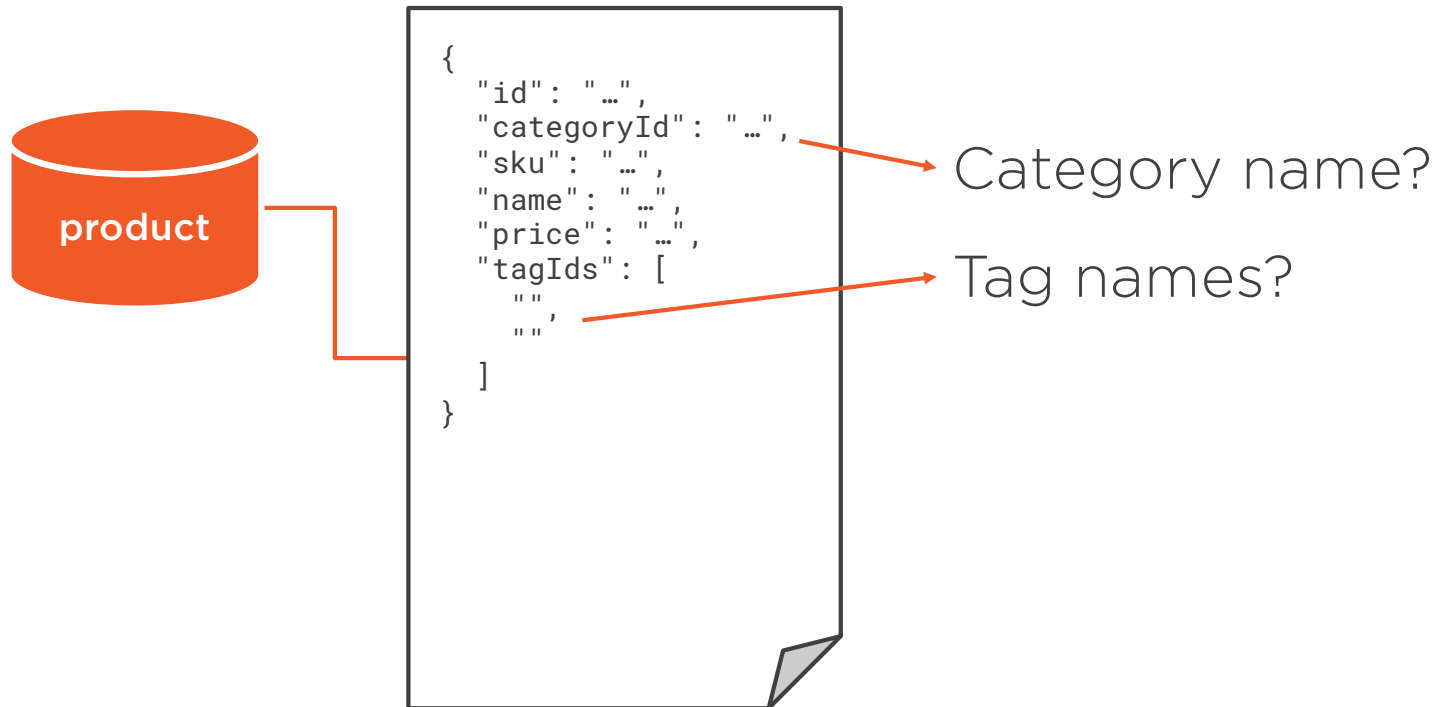
```
SELECT c.name FROM c WHERE c.type = 'category'
AND c.id = '<categoryId>'
```

product
Meta

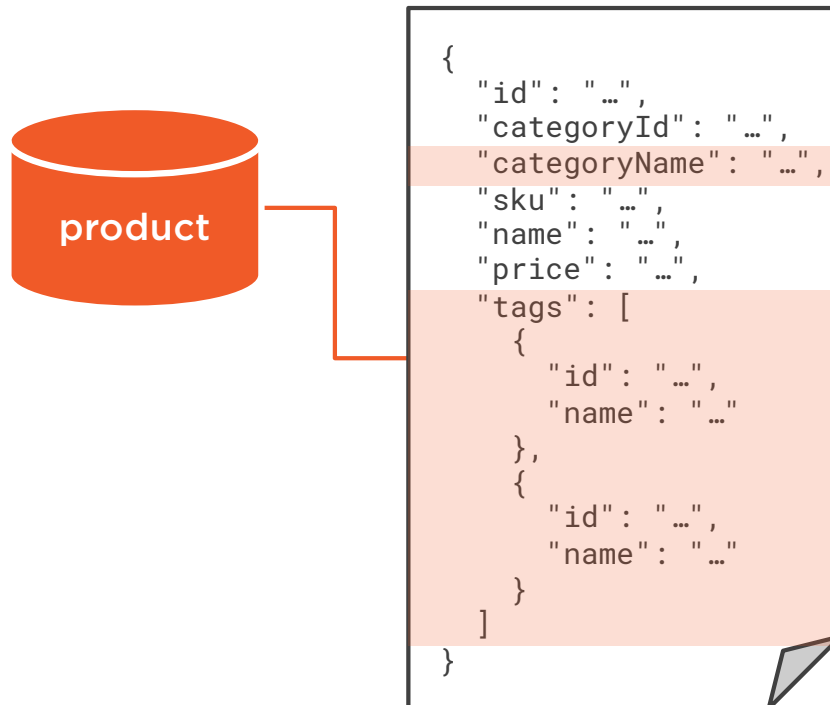
```
SELECT c.name FROM c WHERE c.type = 'tag'
AND c.id IN '<tagId1>, <tagId2>, ..., <tagIdN>'
```



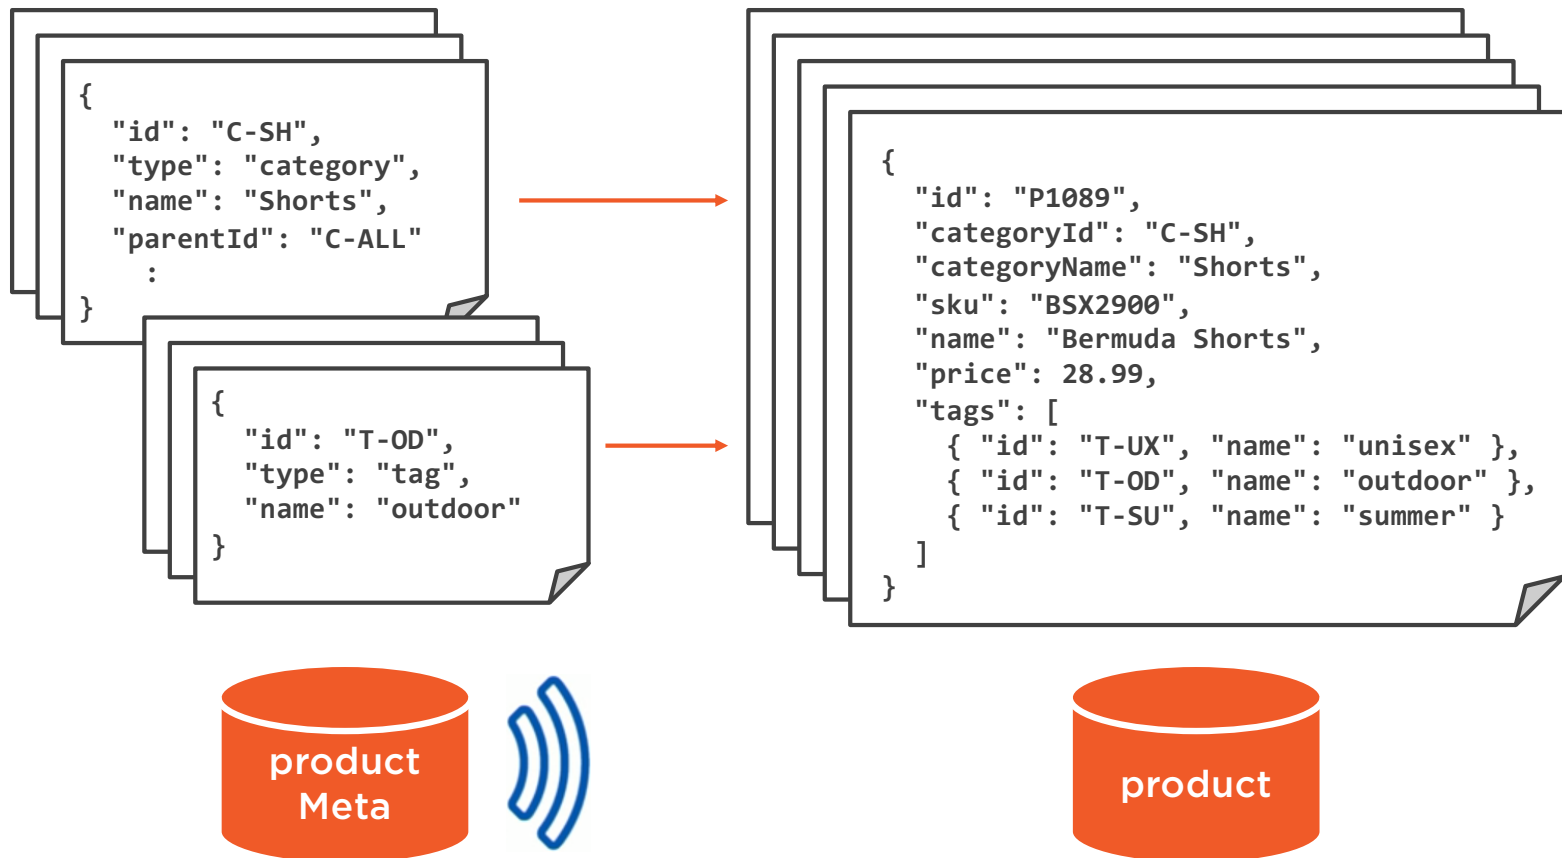
Denormalizing Products



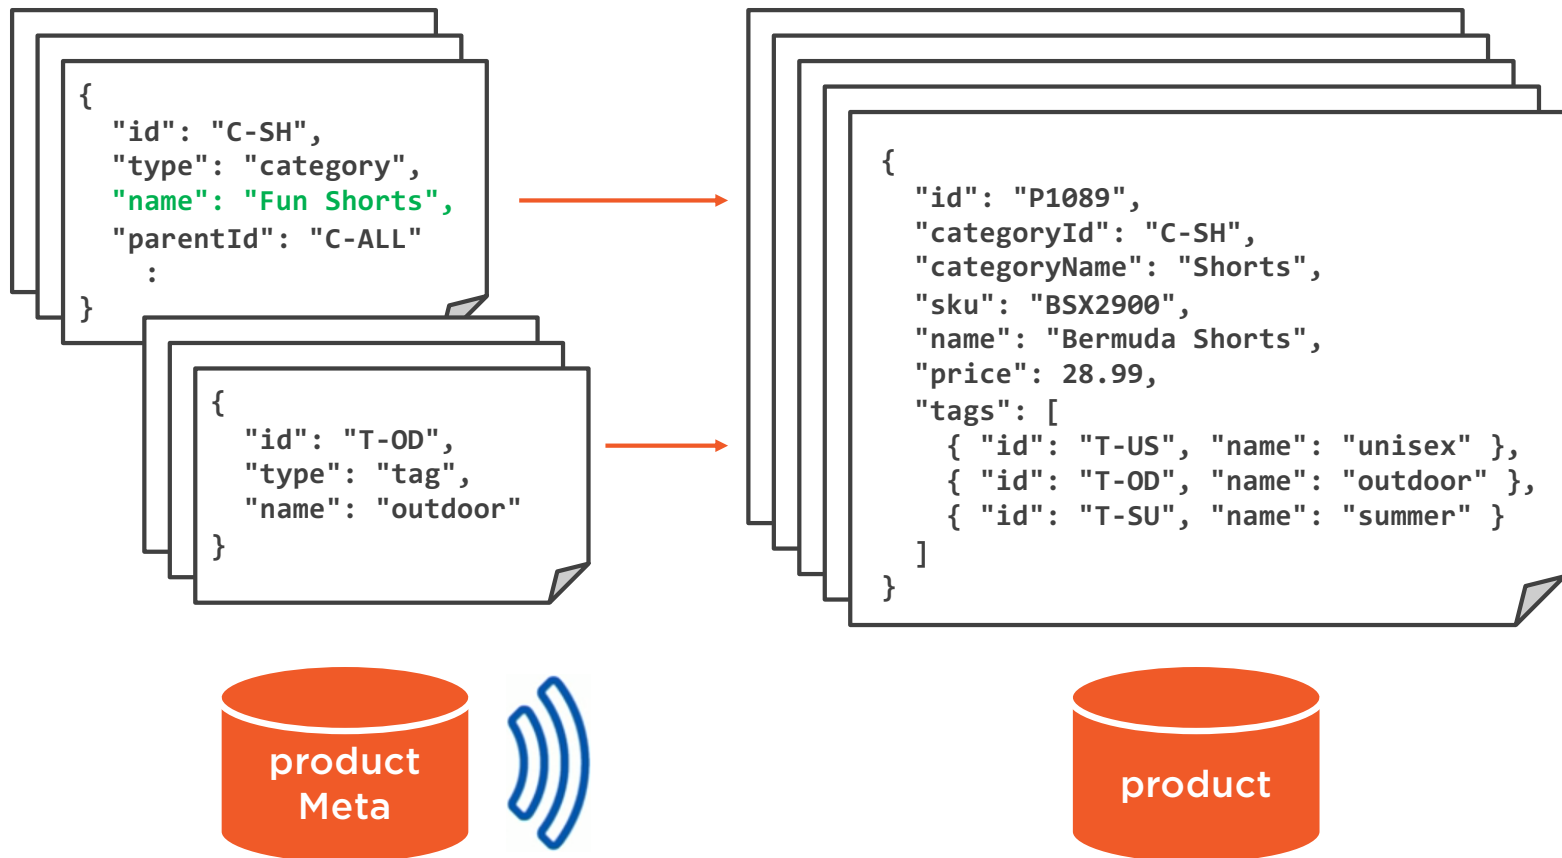
Denormalizing Products



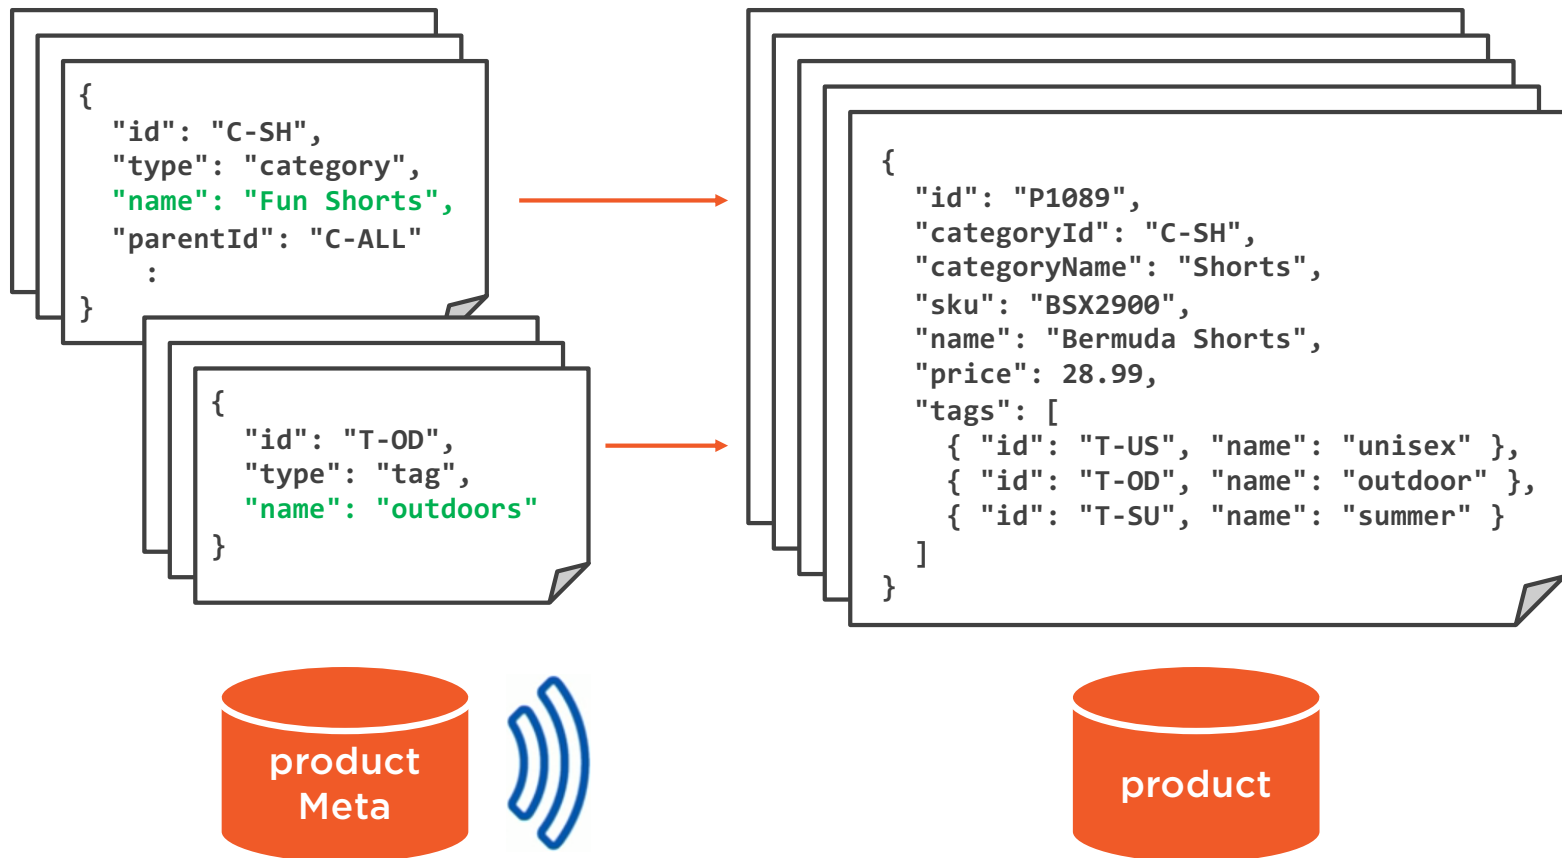
Denormalizing Products



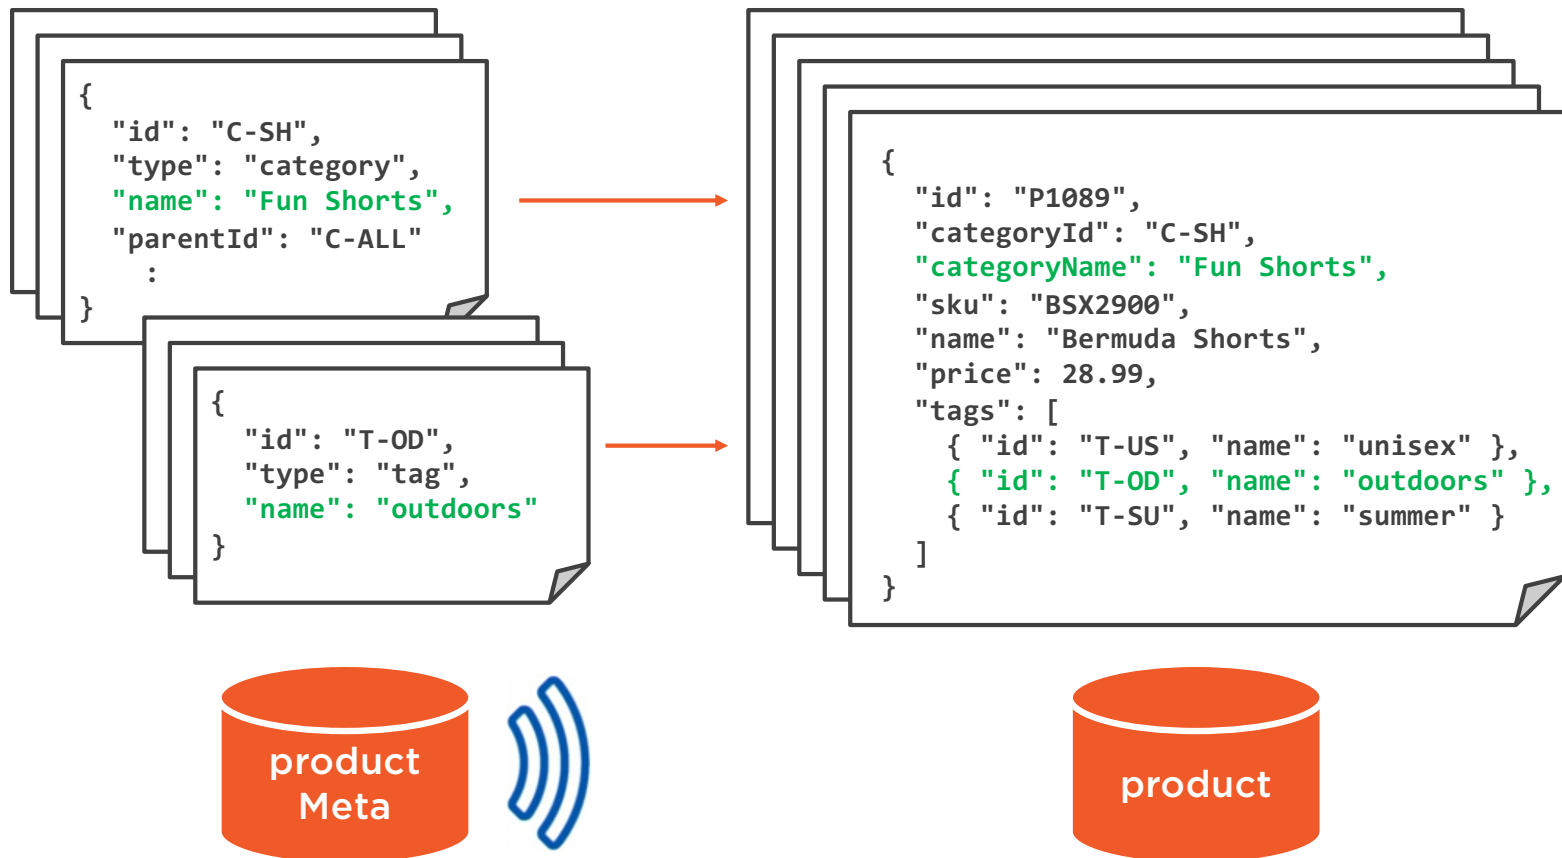
Denormalizing Products



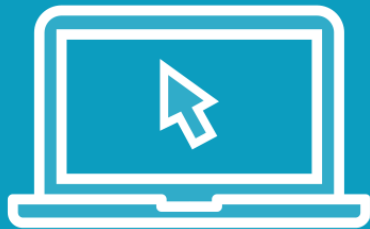
Denormalizing Products



Denormalizing Products



Demo



Denormalization Microservice



Summary



Why denormalize?

- Optimize the data model for common query patterns
- Ensure most queries can be satisfied by a single request

Denormalization microservice

- Product category and tag names

